

# Assignment 1

## Introduction to Simulation with Variance Estimation

Roman Grebnev

2022-10-10

### Task 1

Compare the 4 algorithms against R's 'var' function as a gold standard regarding the quality of their estimates.

- Implement all variants of variance calculation as functions.
- Write a wrapper function which calls the different variants.

Import libraries

```
library("ggplot2")
library("microbenchmark")
```

### Two-pass algorithm

Implementation of **two-pass algorithm**

```
variance_two_pass <- function(vector){
  n = length(vector)
  sample_mean <- sum(vector)/n
  variance <- 1/(n-1) * sum((vector - sample_mean)^2)
  return(variance)
}
```

### Excel algorithm

Implementation of **excel algorithm**

```
variance_excel <- function(vector){
  n = length(vector)
  p1 = sum(vector^2)
  p2 = 1/n * sum(vector)^2
  variance <- (p1 - p2)/(n-1)
  return(variance)
}
```

### Shift algorithm

Implementation of **shift algorithm** with shift by the first element of the vector

```
variance_shift <- function(vector){
  n = length(vector)
  c = vector[1]
  p1 = sum((vector-c)^2)
```

```

p2 = 1/n * (sum(vector-c))^2
variance <- (p1 - p2)/(n-1)
return(variance)
}

```

## Online

Implementation of **online algorithm**

```

variance_online <- function(vector){

  xmn_prev = vector[1] + (vector[2] - vector[1])/2
  varmn_prev = ((vector[2] - vector[1])^2)/2

  for(i in 3:length(vector)) {
    xmn = xmn_prev + (vector[i] - xmn_prev)/i
    varmn = (i - 2)/(i - 1)*varmn_prev + ((vector[i] - xmn_prev)^2)/i

    xmn_prev = xmn
    varmn_prev = varmn
  }

  return(varmn)
}

```

Wrapper function, which produces the result of variance computation for each variance computation algorithm

```

variance_wrapper <- function(vector){
  var_two_pass <- variance_two_pass(vector)
  var_excel <- variance_excel(vector)
  var_shift <- variance_shift(vector)
  var_online <- variance_online(vector)
  return(c(var_two_pass, var_excel, var_shift, var_online))
}

```

## Task 2

Compare the computational performance of the 4 algorithms against R's 'var' function as a gold standard and summarise them in tables and graphically.

Generate vectors with normally distributed values

```

set.seed(12202120)
x1 <- rnorm(100, mean = 0)
set.seed(12202120)
x2 <- rnorm(100, mean=1000000)
set.seed(12202120)
x3 <- rnorm(100, mean=0.000001)

```

Function for benchmark computation

```

micro_benchmark <- function(vect) {
  micro_benchmark_calc = microbenchmark(
    variance_two_pass(vect),
    variance_excel(vect),
    variance_shift(vect),
    variance_online(vect),

```

```

    var(vect)
  )
}

```

Table 1: Comparison of variation computation algorithms for datasets x1 and x2

	two-pass	excel	shift	online
variance_x1	1.254942	1.254942	1.254942	1.254942
variance_x2	1.254942	1.254893	1.254942	1.254942

Table 2: Benchmark summary for generated data with mean = 0

expr	min	lq	mean	median	uq	max	neval
variance_two_pass(vect)	1.250	1.3810	1.51773	1.4655	1.5850	3.386	100
variance_excel(vect)	1.205	1.3445	1.47312	1.4145	1.5440	3.755	100
variance_shift(vect)	1.614	1.8040	2.14572	1.8940	2.0385	21.842	100
variance_online(vect)	9.900	10.0840	10.23886	10.1375	10.2205	15.903	100
var(vect)	5.393	5.6420	6.75919	5.7680	5.9740	97.152	100

Table 3: Benchmark summary for generated data with mean = 1000000

expr	min	lq	mean	median	uq	max	neval
variance_two_pass(vect)	1.078	1.2010	1.36219	1.2965	1.4460	2.817	100
variance_excel(vect)	1.012	1.1575	1.44981	1.2765	1.3810	13.235	100
variance_shift(vect)	1.436	1.5810	1.76087	1.6535	1.8495	4.678	100
variance_online(vect)	8.829	9.0190	9.80795	9.1590	9.6415	16.055	100
var(vect)	4.684	4.9155	5.43554	5.0645	5.3795	21.369	100


```

knitr::opts_chunk$set(fig.width=20, fig.height=8)
par(mar=c(5, 12, 5, 5))
boxplot(log(time)~expr,
        data=benchmark_x1,
        main="Benchmarks of algorithms' execution speed,
              mean = 0",
        xlab="Log(Time)",
        col="red",
        border="blue",
        horizontal=TRUE,
        las=2
)

par(mar=c(5, 12, 5, 5))
boxplot(log(time)~expr,
        data=benchmark_x2,
        main="Benchmarks of algorithms' execution speed,
              mean = 1000000",
        xlab="Log(Time)",

```

```
col="red",
border="blue",
horizontal=TRUE,
las=2
)
```



## Conclusion

### Summary for comparison of variance computation algorithms

- Variance computed with all algorithms is the same for dataset x1 with mean = 0
- Variance computed for dataset x2 with mean = 1000000 is the same for algorithms, but results produced by excel algorithm, which has differences starting from 4 decimal place

### Summary for benchmarking performed for x1 dataset with mean = 0 and x2 dataset with mean = 1000000

- The ranking of the performed benchmarking is the same for both datasets
- Speed of variance computation for all algorithms is slightly higher for dataset x1 with mean = 0 than for x2
- According to the performed benchmarking, **excel** and **two-pass** algorithms have the highest execution speed.
- Gold-standard implementation of **var** function in R is on the 4th place, while **excel** and **two-pass** outperform it in terms of computation speed
- The least performant algorithm is **online** implementation, because of the iterative approach to updating the variance with each new element of vector

## Task 3

Investigate the scale invariance property for different values and argue why the mean is performing best as mentioned with the condition number.

- Compare the results according to the instructions provided by Comparison I and Comparison II of the slides.
- Provide your results in table format and graphical format comparable to the example in the slides.

*Description of the approach* To investigate the property of scale invariance we can perform the following experiment: - Generate the sequence of values within the range between the minimum and maximum values of distributions - Iteratively compute the condition number for each obtained shift constant - Compare the results and identify, whether condition number is affected by the shift constant  $c$

**Shift implementation** of variance algorithm

```

func_cond_num <- function(vector, const){
  n = length(vector)
  vector_shifted = vector - const

  mean_shifted <- mean(vector_shifted)

  s_shifted = sum((vector_shifted - mean_shifted)^2)
  cond_number <- sqrt(1 + (mean_shifted^2 * n)/s_shifted)

  return(cond_number)
}

func_compare_cond_numbers <- function(vector){
  cond_numbs <- c()

  consts <- seq(from = min(vector), to = max(vector), length.out = 100) #((max(vector)-min(vector))/100)
  consts <- append(consts, mean(vector))
  consts <- sort(consts)
  for(const in consts) {
    cond_number <- func_cond_num(vector, const)
    cond_numbs <- append(cond_numbs, cond_number)
  }
  df <- data.frame(cond_numbs, consts)
  return(df)
}


```

```

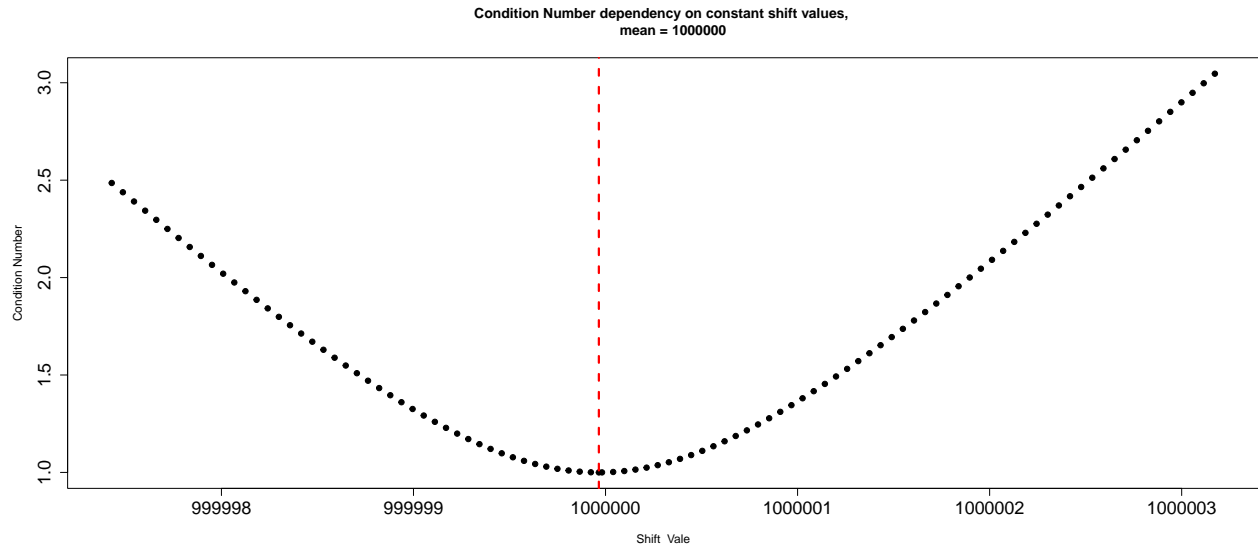
knitr::opts_chunk$set(fig.width=20, fig.height=12)

df_cond_number_x1 <- func_compare_cond_numbers(x1)
par(mar=c(5, 12, 5, 5))
plot(df_cond_number_x1$consts, df_cond_number_x1$cond_numbs, main="Condition Number dependency on constant shift values,
      mean = 0",
      xlab="Shift Value", ylab="Condition Number", pch=19, cex.axis = 1.5)
abline(v = mean(x1), col="red", lwd=3, lty=2)

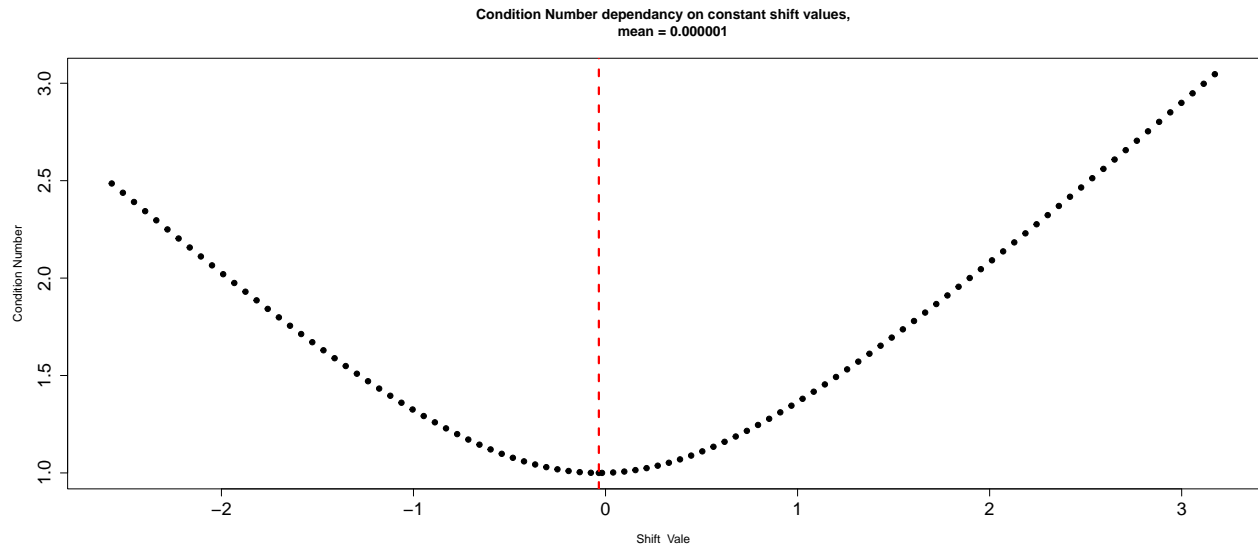
```



```
df_cond_number_x2 <- func_compare_cond_numbers(x2)
par(mar=c(5, 12, 5, 5))
plot(df_cond_number_x2$consts, df_cond_number_x2$cond_numbs, main="Condition Number dependency on constant shift values, mean = 1000000",
     mean = 1000000,
     xlab="Shift Vale", ylab="Condition Number", pch=19, cex.axis = 1.5)
abline(v = mean(x2), col="red", lwd=3, lty=2)
```



```
df_cond_number_x3 <- func_compare_cond_numbers(x3)
par(mar=c(5, 12, 5, 5))
plot(df_cond_number_x3$consts, df_cond_number_x3$cond_numbs, main="Condition Number dependency on constant shift values, mean = 0.000001",
     mean = 0.000001,
     xlab="Shift Vale", ylab="Condition Number", pch=19, cex.axis = 1.5)
abline(v = mean(x3), col="red", lwd=3, lty=2)
```



## Conclusion

The vertical dashed line highlights the mean of the distribution. The lowest condition number is observed at the dotted line, which indicates that the optimal condition number  $k \sim 1$  is obtained when the shift constant equals the mean of the distribution. It leads to the conclusion that the best condition number  $k = 1$  is

achieved for the constant shift of distribution by the mean of the distribution. Change of the shift constant  $c$  has the same effect on the condition number for all datasets (with mean = 0, mean = 1000000, mean = 0.000001)

### Comparison of variance computation for different values of constant shift $c$

Description of the approach: In order to demonstrate the impact of constant  $c$  on the robustness of the variance computation, we can introduce noise into original dataset  $x_1$  with mean = 0. We can iterate over the range of shift constants  $c$  from  $\min(\text{vector}) * 1000000$  to  $\max(\text{vector}) * 1000000$  and compare the deltas between the variance, computed for original data and for noisy data.

```
func_var_const <- function(vector, const){
  n = length(vector)
  vector_shifted = vector - const

  mean_shifted <- mean(vector_shifted)

  variance <- sum((vector_shifted - mean_shifted)^2)/(n-1)

  return(variance)
}

variance_shift_c <- function(vector, const){
  n = length(vector)
  p1 = sum((vector-const)^2)
  p2 = 1/n * (sum(vector-const))^2
  variance <- (p1 - p2)/(n-1)
  return(variance)
}

compute_noisy_variance <- function(vector){

  noise <- rnorm(100, mean = 0) * 0.001
  constant_shifts <- seq(from = min(vector)*1000000, to = max(vector)*1000000, length.out = 1000)

  variances <- c()


  vector_n <- vector + noise

  for(const in constant_shifts) {
    variance <- variance_shift_c(vector_n, const) - variance_shift_c(vector, const)
    #variance_diff <- func_var_const(vector_n, const) -
    variances <- append(variances, variance)
  }

  df <- data.frame(variances, constant_shifts)
  return(df)
}

df_noisy <- compute_noisy_variance(x1)

plot(df_noisy$constant_shifts, df_noisy$variances, xlab = "Constant shift", ylab = "Delta of variance b
abline(v = 0, col="red", lwd=3, lty=2)
```



### Summary of shift constant experiment with added noise

Implemented procedure indicates, that variance deltas between original and noisy data has the biggest variation with shift constant  $c$ , far from distribution mean, on the contrary computation becomes more robust, when shift constant  $c$  approaches the distribution's mean. This illustrates, that shift constant  $c$  affects condition number, which in turn creates the effect of noise amplification.

### Why shift by $c = \text{mean}(\text{data})$ has the best condition number $k$ ?

With  $k$  computed as follows:  $\bar{k} = \sqrt{1 + \frac{n}{S} (\bar{x} - c)^2}$  It means that the most robust variance computation, when  $k = 1$  can be obtained with the shift  $c = \bar{x}$ .

### Task 4

Compare condition numbers for the 2 simulated data sets and a third one where the requirement is not fulfilled, as described during the lecture.

```
df_cond_number <- data.frame(rbind(func_cond_num(x1, 0), func_cond_num(x2, 0), func_cond_num(x3, 0)))
colnames(df_cond_number) <- c("Condition number")
rownames(df_cond_number) <- c("x1 (mean = 0)", "x2 (mean = 1000000)", "x3 (mean = 0.000001)")
kable(df_cond_number, caption = "Condition numbers for datasets x1, x2, x3")
```

Table 4: Condition numbers for datasets x1, x2, x3

	Condition number
x1 (mean = 0)	1.000485e+00
x2 (mean = 1000000)	8.971613e+05
x3 (mean = 0.000001)	1.000485e+00



## Conclusion

- Condition numbers for datasets x1 and x3 are the same
- For dataset x2 condition number is very high, which means that variance is less robust to the errors. There is a practical sense to perform shift corrections by sample mean to reduce the condition number  $k$ , which will result in more robust variance computation.